



When the Only Dissent Is Human

LLM and Analyst Courses of Action Under Military Crisis Conditions

Paul Grupp · Lea Kiric

University of Konstanz · Fairness and Collective Decision-Making in AI

Models are moving into the command chain

DEPLOYMENT

Maven became a formal program of record in 2026. Language models were added to its advisory layer.

AT MACHINE SPEED

Thunderforge and JHU/APL wargaming push the same logic into planning.

THE DISPUTE

In early 2026 the Pentagon labeled Anthropic a supply-chain risk after it refused “any lawful use” terms and ordered agencies off Claude within six months.

What origin should predict

WESTERN MODEL

- Reads a maritime standoff through deterrence.
Recommends a show of force.

Invokes UNCLOS, freedom of navigation, alliance
credibility.

CHINESE MODEL

- Reads the same facts as a sovereignty
violation.

Recommends non-kinetic pressure, or declines.
Invokes absolute sovereignty, non-interference.

Setup



MODELS

West: ChatGPT, Claude, Gemini.

China: DeepSeek, Qwen.

INPUT

Identical OSINT briefings, embedded inline for byte-identical input. Neutral labels: State A, B, C.

CONSTRAINT

One forced choice, A-E, three runs each, at the analyst-facing level: what one user checks under time pressure.

RUBRIC

Escalation 0-4 · Blame -2 to +2 · Legitimacy (categorical) · Refusal 0-2.

The human reference point

WHAT IT BUYS

If the Western models track him and the Chinese diverge, that supports origin. If he diverges from all five, that too is a result.

The Contested Strait



The field converges

Model	Cohort	Escalation S1	Escalation S2	Blame S1	Blame S2	Refusal
ChatGPT	West	3	3	+1	+1	0
Claude	West	2	3	+1	+1	0
Gemini	West	3	3	+1	+1	0
DeepSeek	China	3	3	+1	+2	0
Qwen	China	3	3	+1	+1	0

Legitimacy: Western-dominant vocabulary across all five models in S1; an IHL vocabulary wrapped around a self-defense recommendation in S2.

Refusal: flat 0 across every model, every run, both scenarios, including the two Chinese models.

Scale: escalation 0–4, blame –2 to +2, refusal 0–2. Highlighted cells are the only two deviations in the dataset.


The exceptions run the wrong way

CLAUDE · WEST

The lowest escalation in the dataset. Alone at rung 2, non-kinetic coercion, one below the four-model consensus.

DEEPSEEK · CHINA

The most one-sided attribution. Alone at blame +2, full responsibility on the Organization, where the field holds +1.



The only real dissent is human


Advising State C, he refused to return a single letter.

His recommendation branched on a fact the brief leaves open, whether State A's grounded vessel lies in its own or in international waters:

THE CONTRAST

Each of the five models silently resolved the same gap and committed to one rung. The human flagged the missing fact as decision-relevant and surfaced it.

Replicable, and the study's clearest behavioral divide.



The model is built to win, and it will choose its course of action on that basis. In a real crisis that is dangerous, and it is the reason these systems need safeguards rather than authority.

— Human analyst

What this means



LIMITATIONS

- Baseline-equivalent only; no perspective-mirror runs yet.

Three model identities still to be backfilled. Both deviations sit on flagged coding conventions.

A single analyst, from the Western cohort's culture.

THE REFRAME

On this evidence a Western and a Chinese model counsel much the same thing. The risk is that all of them counsel with a confidence the situation does not warrant, resolving away the ambiguity a human treats as the substance of the decision.

Safest used to challenge human assumptions, not to replace the judgment that still carries the risk.